

Bing Tea Search Engine



Jimmy Dai:	CSE Junior
Rickey Dong:	CSE Junior
Kevin Li:	CSE Senior
Roe Tsimhoni:	CSE Junior
Matthew Wang:	CSE + Aero Junior
Danny Kim:	CS Junior
Wonbin Jin:	CSE Junior



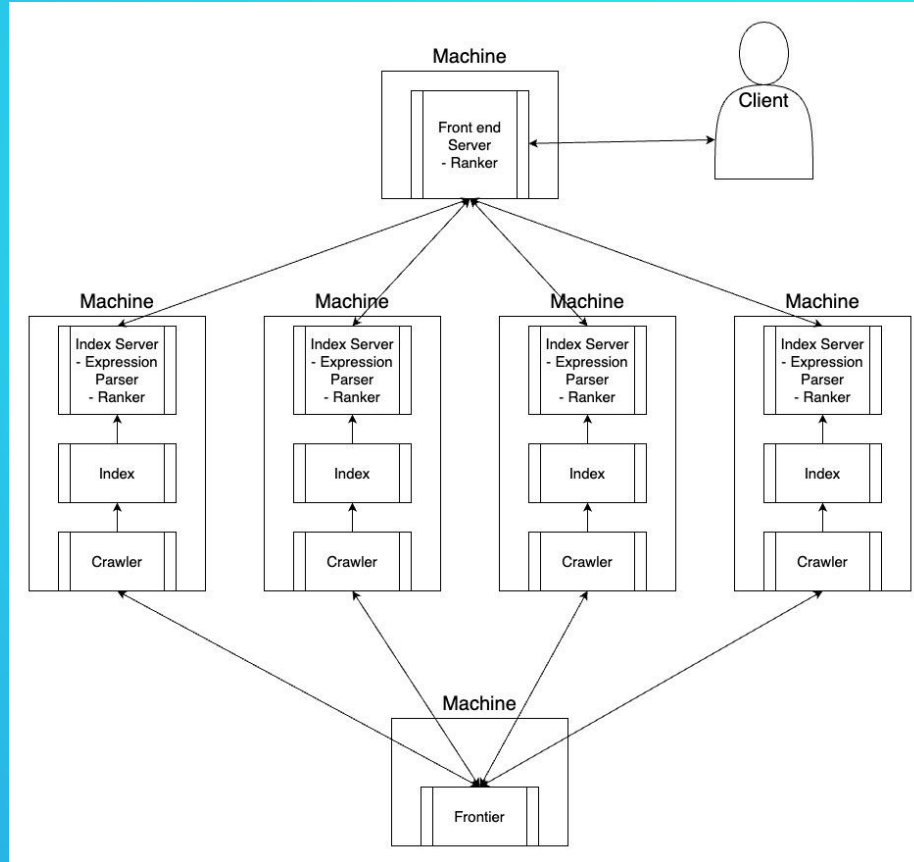
0:00



BING TEA

Search Bing Tea...

Architecture



Architecture

Key Design Choices:

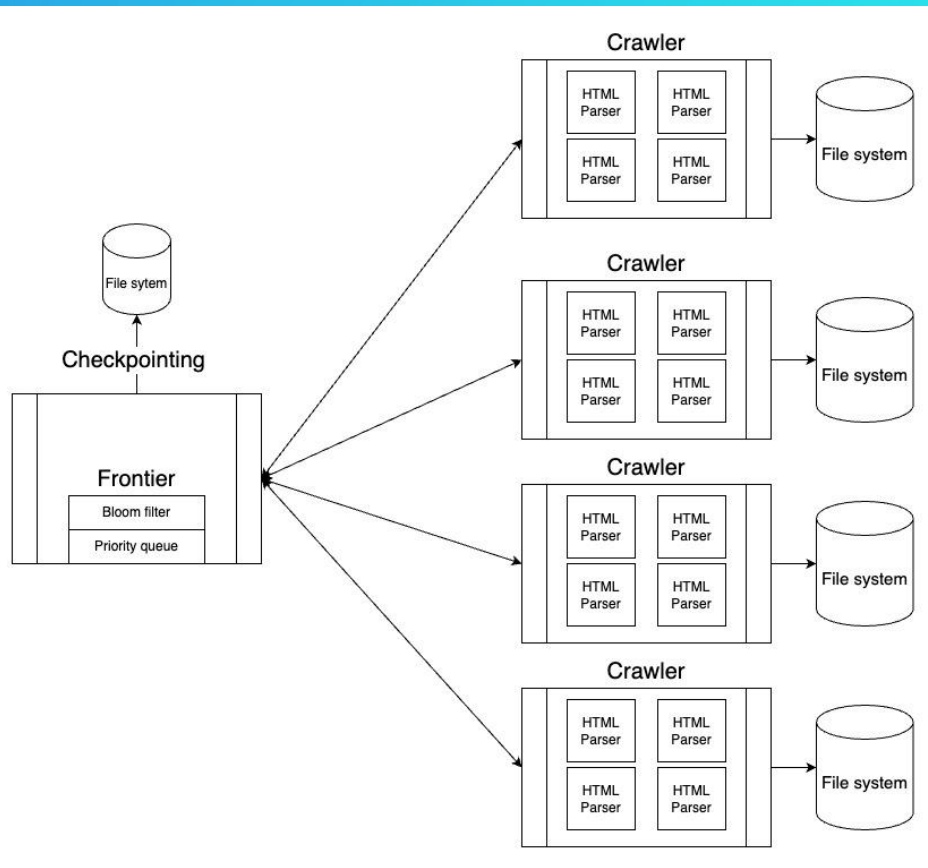
- A “manager” process for the crawler
- No “special” indexes
- Store parsed HTML files
- Closeness from seed-list static ranking
- Two stage ranking

Crawler

Statistics:

Total Pages Crawled	51,618,383
Total Crawl Time	170 hours
Average documents per second	84 documents / second
Peek documents per second	215 documents / second
Average Frontier response time	2.2 ms

Crawler



- Initial assumption was to crawl on laptops (x7)
- Used 22 VMs on GCP instead
- Average Frontier message wait time: 20ms
- Average Frontier message process time: 2ms
- Easy to add and remove machines **during** crawl

Crawler

Problems and Fixes/Potential Improvements:


- Being stuck in a certain host → Random url crawling, inversely rank hosts by crawl count
- Document retrieval success rate is low → Retried depending on error, figure out a better retrieval method, rank hosts by success rate
- Duplicate detection → Bloom filter on Frontier, bloom filter on Crawler
- Client-side dynamic content → JS rendering or smart heuristics to discover links behind dynamic content
- Page type reach gaps → Whitelist pdf and video types and extract pdf text & video metadata

HTML Parser

- Parses input html and returns key metrics for ranking and index building
- Collects the following data:
 - Title Words
 - URLs and Anchor text
 - “Emphasized” words (italics, bold, headers)
 - Normal words
 - Number of images
 - Site language
- Does some basic text processing for the index
 - Lowercases all words
 - Removes unnecessary punctuation

Bridge Between Parsing and Indexing

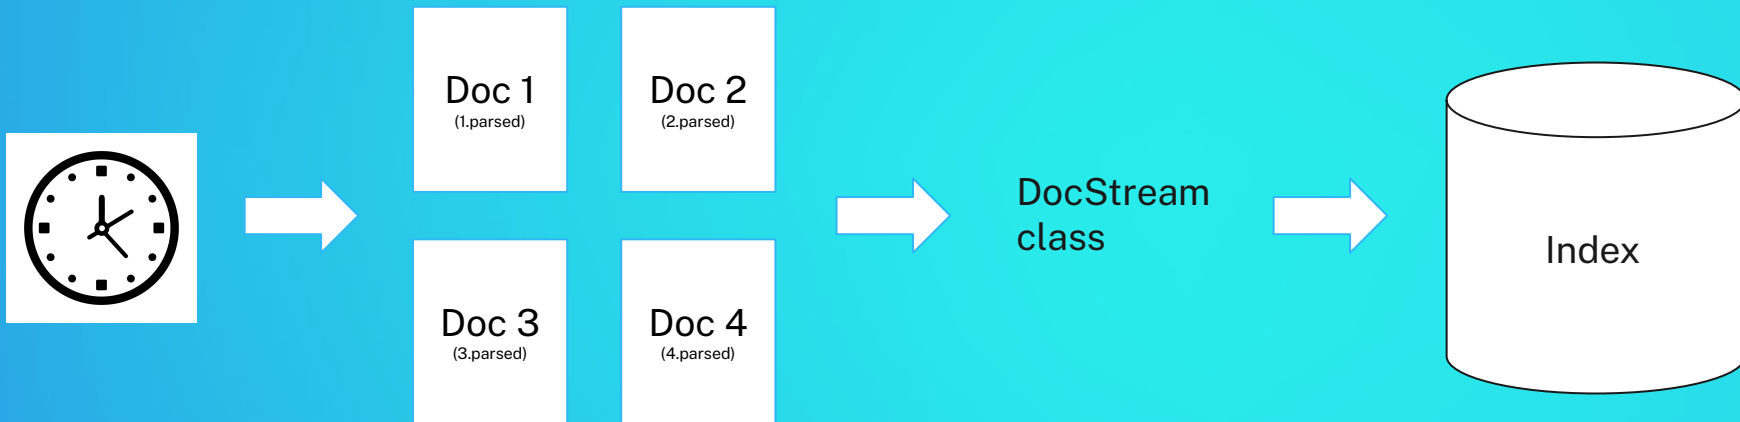
One major design decision we made was to have an explicit phase / “bridge” between parsing and indexing



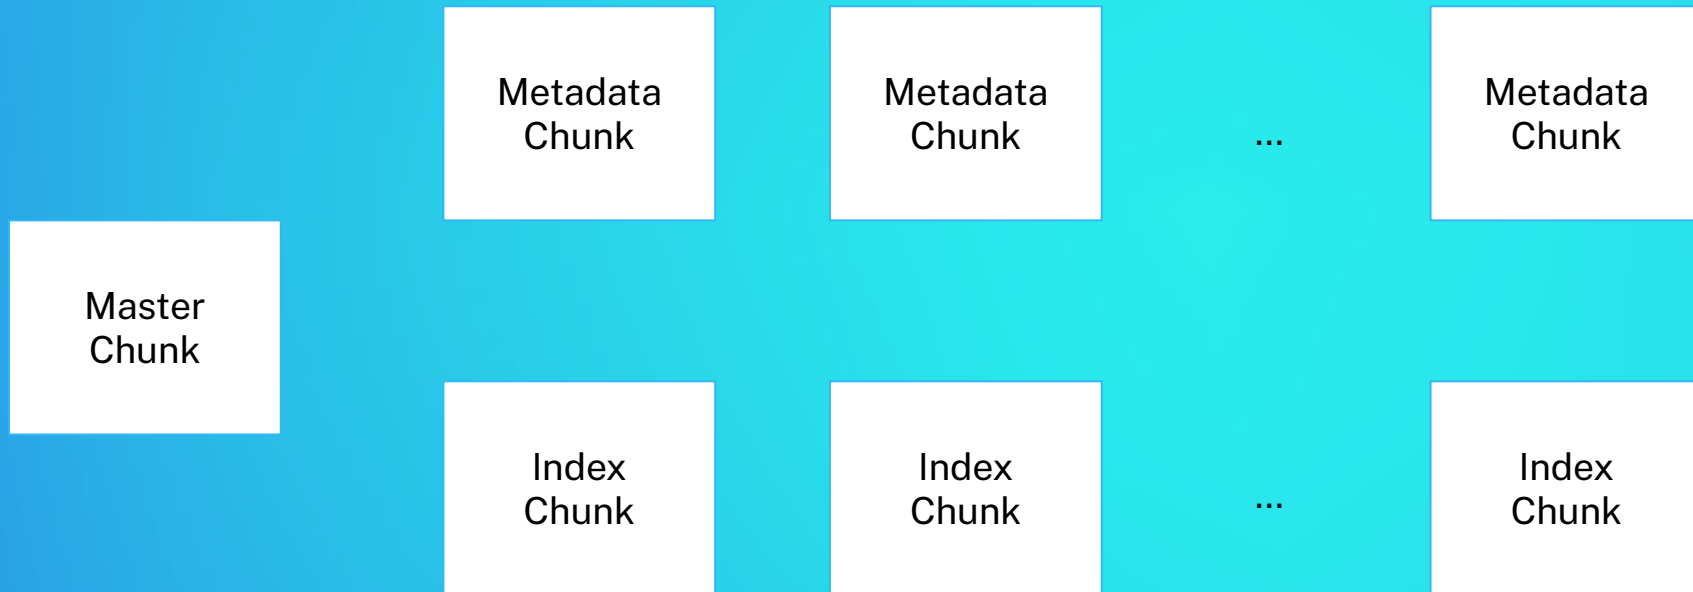
Reasons why this decision was made:

- We wanted to get crawling done ASAP but the index was not yet finished
- We didn't want the index crashing to affect the crawl
- These intermediate parsed files turned out to be good for snippets

Index



Index



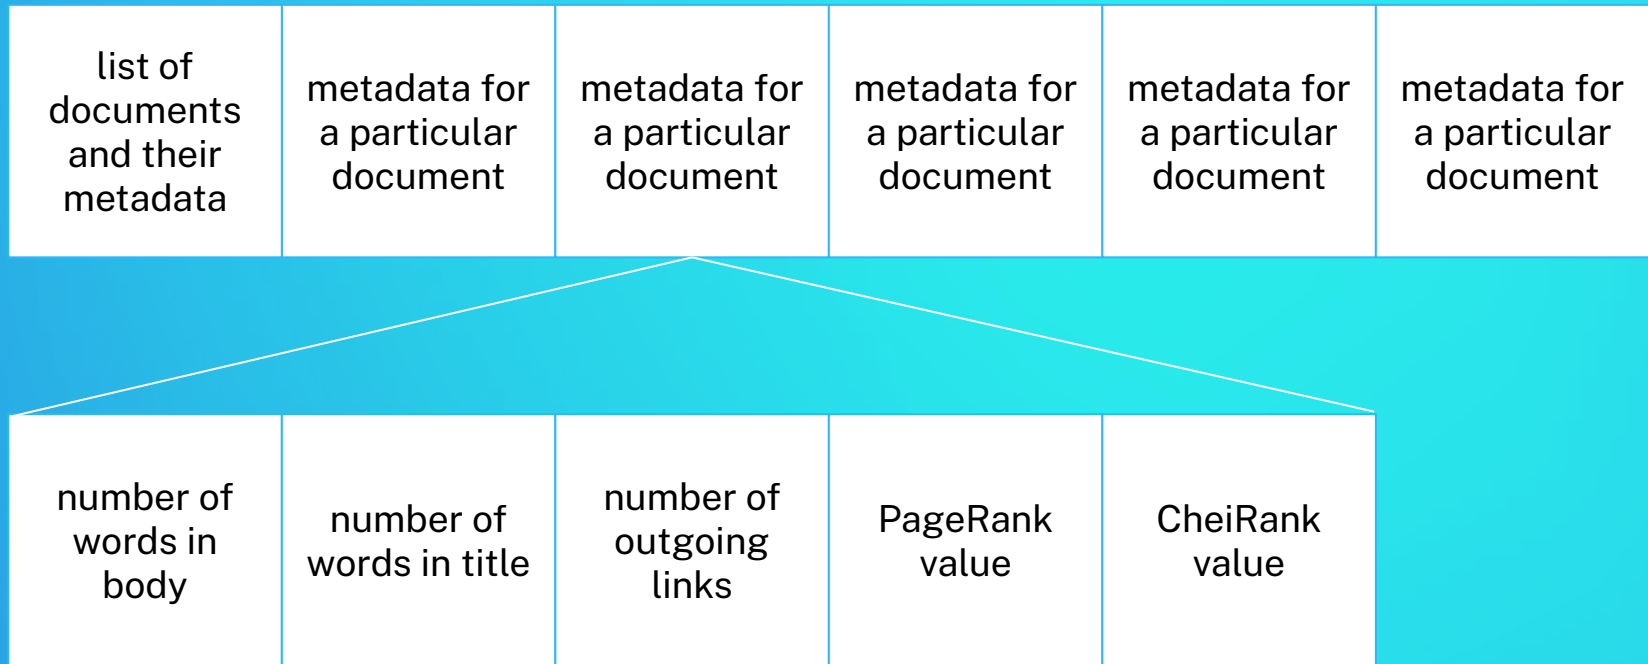
Index

Master Chunk

metadata like # of documents	list of index chunks	list of metadata chunks
------------------------------------	-------------------------	-------------------------------

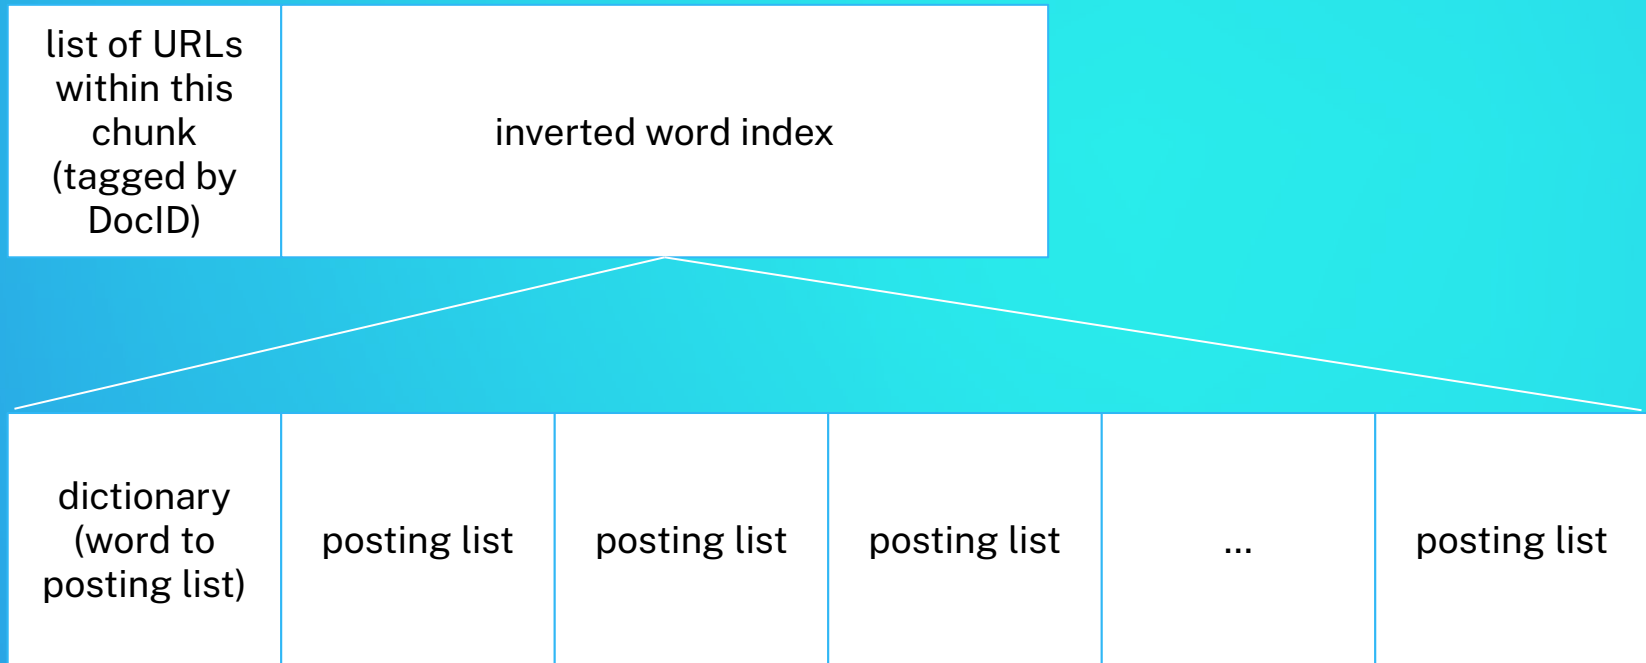
Index

Metadata Chunk



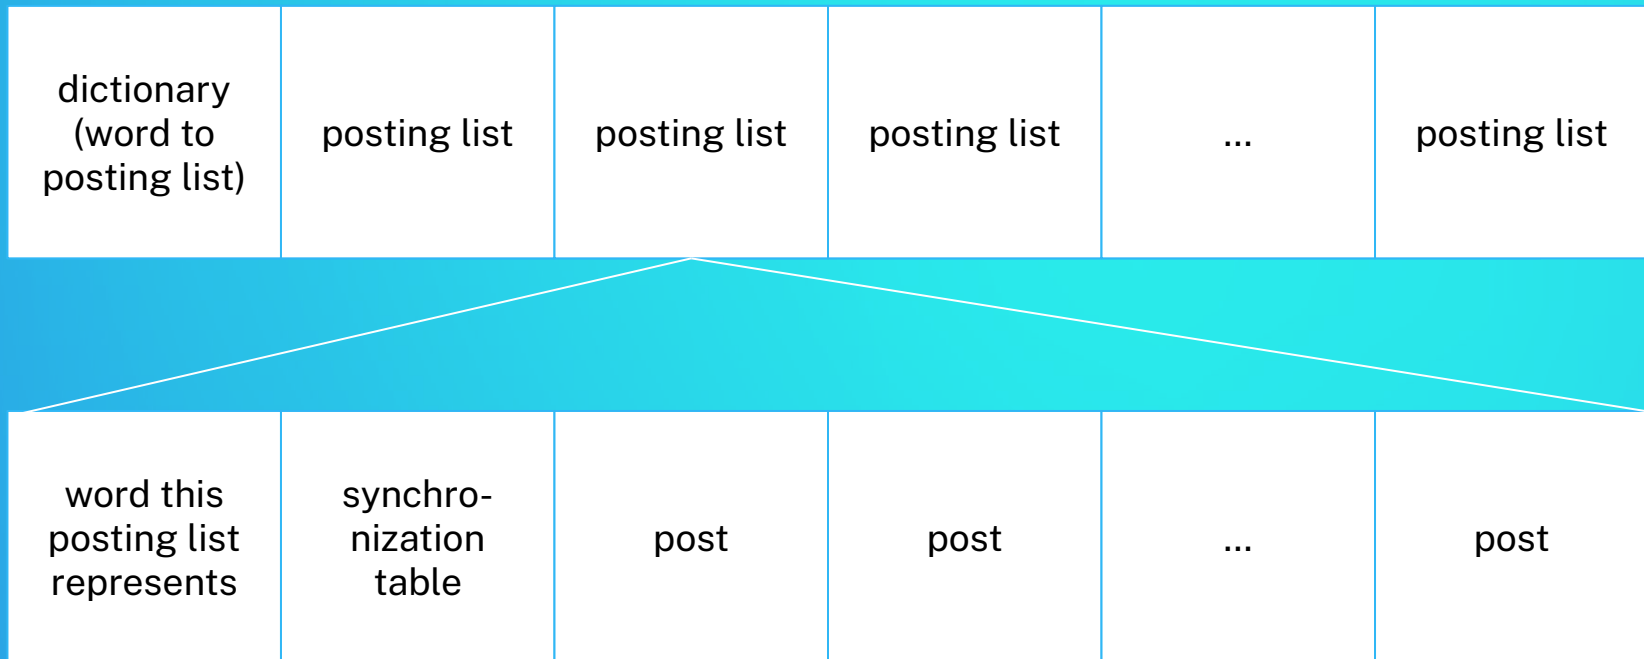
Index

Index Chunk



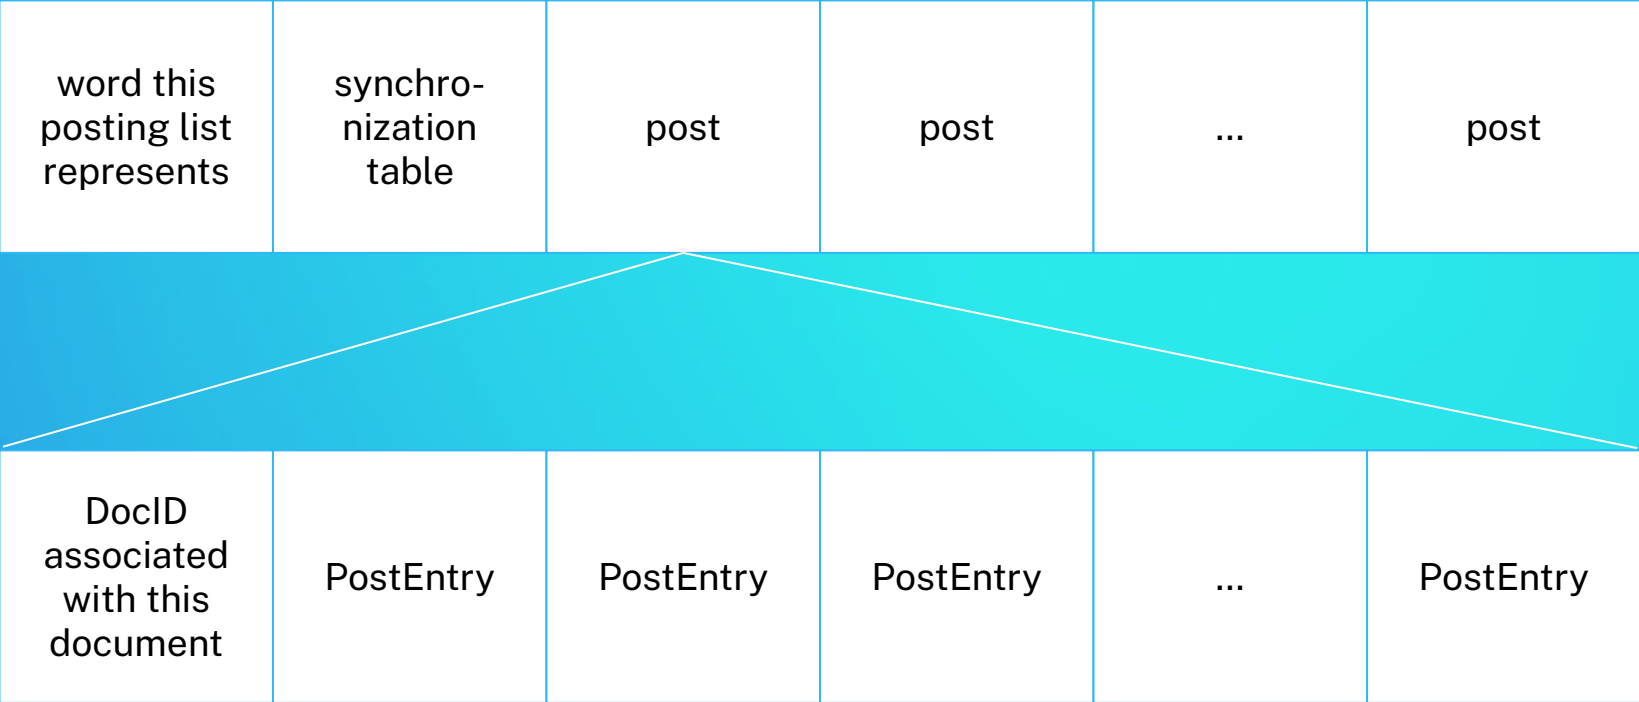
Index

Posting List



Index

Post



Index

PostEntry

DocID associated with this document	PostEntry	PostEntry	PostEntry	...	PostEntry
--	-----------	-----------	-----------	-----	-----------

relative delta from previous PostEntry	specific data about this occurrence (like bold/ital)
---	--

Compression and Synchronization

- postings lists are compressed with variable byte encoding
- instead of storing absolute locations, relative offsets from the previous entry are stored
- these relative offsets are compressed with vByte encoding (not UTF-8)

docIDs	824	829	215406
gaps		5	214577
VB code	00000110 10111000	10000101	00001101 00001100 10110001

- for the ISR functions that Seek() to a particular location, or need to “jump ahead” at a particular spot, instead of moving forward the ISRs one by one an inch at a time...
- we (will) use synchronization points to help scrub the ISRs forward much faster

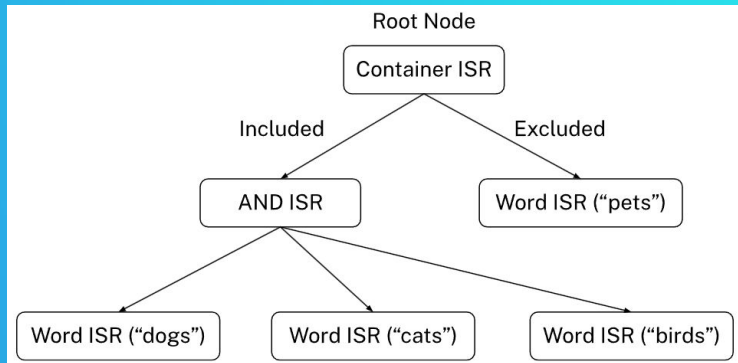
Query Compiler

- Parses and tokenizes input query
- Uses Top Down Recursive Descent (TDRD) to construct an ISR tree structure
- Supports the following BNF Grammar:
 - **OR:** $\langle \text{Constraint} \rangle ::= \langle \text{BaseConstraint} \rangle \{ \langle \text{OrOp} \rangle \langle \text{BaseConstraint} \rangle \}$
 - **AND/NOT:** $\langle \text{BaseConstraint} \rangle ::= \langle \text{SimpleConstraint} \rangle \{ [\langle \text{AndOp} \rangle] \langle \text{SimpleConstraint} \rangle \} | \langle \text{SimpleConstraint} \rangle \langle \text{NotOp} \rangle \langle \text{SimpleConstraint} \rangle$
 - **PHRASE:** $\langle \text{Phrase} \rangle ::= \text{' ' } \{ \langle \text{SearchWord} \rangle \} \text{' '}$
 - **NESTED:** $\langle \text{NestedConstraint} \rangle ::= \text{' (} \langle \text{Constraint} \rangle \text{')'}$
 - $\langle \text{SimpleConstraint} \rangle ::= \langle \text{Phrase} \rangle | \langle \text{NestedConstraint} \rangle | \langle \text{SearchWord} \rangle$
- WordISR retrieves posting lists from the index
- Other ISRs (Not/Container, And, Or, Phrase) operate on children ISRs

Query Compiler

Example Query: (Dogs Cats Birds) NOT Pets

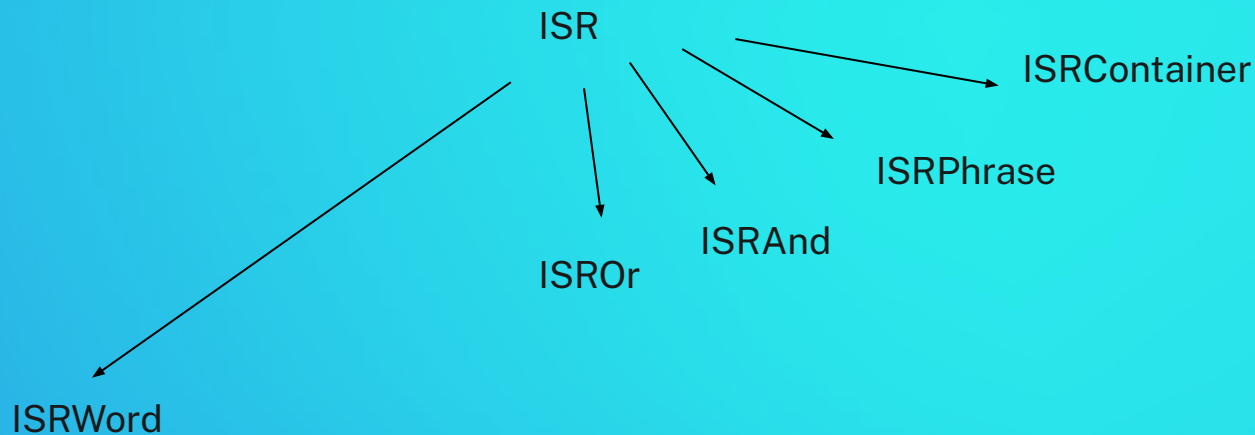
1. Parse and Tokenize:
LPAREN, Word("dogs"), Word("cats"), Word("birds"), RPAREN, NOT, Word("pets")
2. Construct expression tree using BNF language
3. Evaluate expression tree and construct ISR tree structure



4. Search using the Root Node of the ISR tree

Index Stream Readers

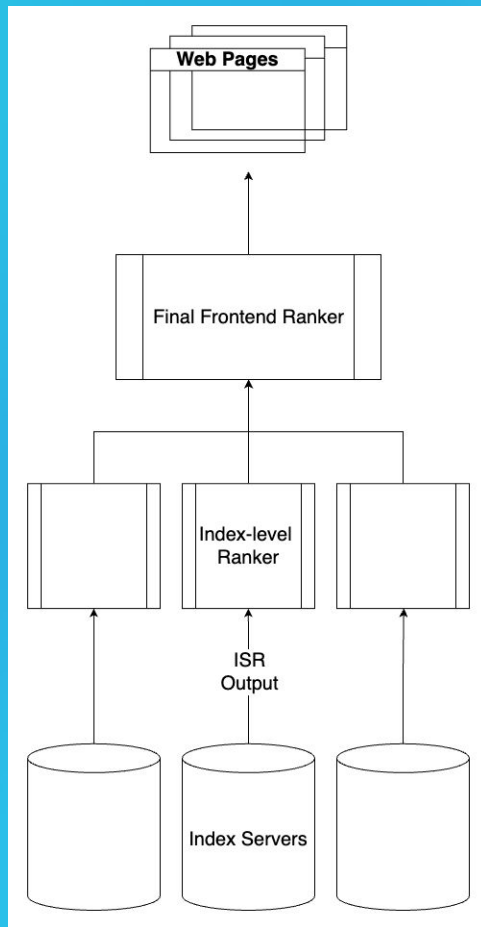
- provides an interface to read through the index and find occurrences of a particular term
- Next(), NextDocument(), Seek()



Ranker

Architecture

- ML-based ranker using XGBoost
 - Data (Pages) described through heuristic metadata scores
 - Ranker selects outputs (pages) with the top 10 ranks
- 2-stage ranking system
 1. Index-level ranker (Heuristic score)
 2. Final Frontend ranker (Predicted score)



Ranker

Heuristics

Pages heuristically “scored” based on page metadata

Static Ranking:

- Document Length
- Title Length
- Page Rank
- CheiRank
- Number of Images
- Number of Outgoing Links
- Domain


Dynamic Ranking:

- Weighted sum score based on query match locations
 - Title matches have higher weight compared to body matches

Ranker

PageRank, CheiRank, and Louvain

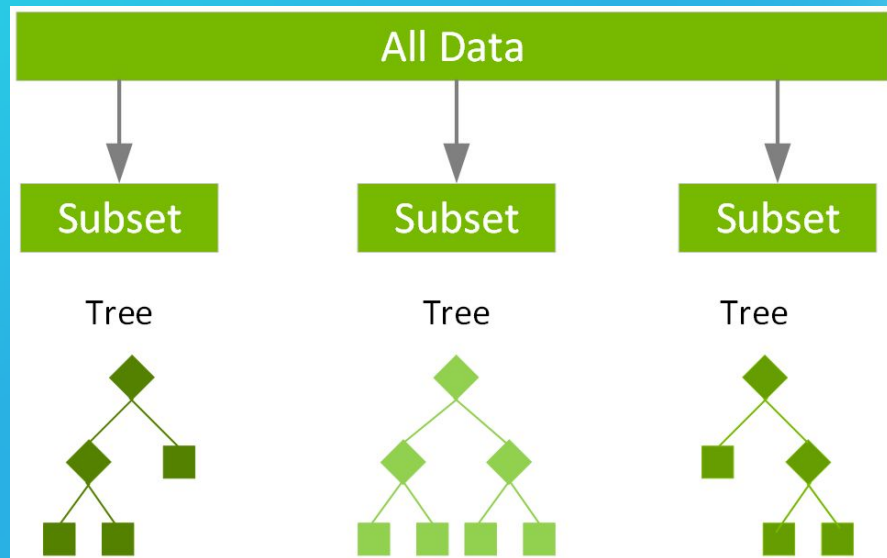
- Network-based
- Consider link structure and between-site relationships.
- PageRank highlights popular nodes, while CheiRank highlights communicative nodes.
- Louvain detects communities by greedily maximizing modularity.
- Communities likely share topics and have similar goodness.



Ranker

XGBoost

- Instead of a linear model, we use a more advanced model to extract more nuance.
- eXtreme Gradient Boosting (XGBoost) is a gradient boosted decision tree model.
- It trains quickly and accurately, and is commonly used in data science competitions.
- Often needs less data and faster to train than neural networks, while performing better.



Frontend

 **ING TEA**

Search Bing Tea...

 **ING TEA**

"university of michigan"

Computer Science Major Eng Computer Science and Engineering at Michigan

<https://cse.engin.umich.edu/academics/undergraduate/programs/computer-science-eng/>

Due to capacity constraints students who are admitted to the University of Michigan in Fall 2023 or later must first

Computer Science and Engineering at Michigan

<https://cse.engin.umich.edu/>

We are part of the College of Engineering at the University of Michigan 2025 The Regents of the University

Brian D Noble 8211 Home of Brian D Noble

<https://bnoble.engin.umich.edu/>

and Computer Science in the College of Engineering at the University of Michigan He is well known as an advocate

Andrew DeOrio

<https://andrewdeorio.com/>

teaching faculty and Associate Chair for Undergraduate Affairs at the University of Michigan and a consultant for web projects His

University of Michigan

<https://umich.edu/>

University of Michigan

Nicole Hamilton

<https://nicolehamilton.com/>

Nicole Hamilton

Nicole Hamilton

<https://web.eecs.umich.edu/~nham/>

Personal Lecturer III 2017 to 2021 Computer Science and Engineering University of Michigan Ann Arbor MI 48109 2121 Cell nbsp